

# CP386: Assignment 3 – Winter 2023

**Due on March 3, 2023 (Before 11:59 PM)**

This is a group (of two) assignment, and we will practice the concept of the multi-threading, process scheduling, and some related system calls.

## General Instructions:

- For this assignment, you must use C language syntax. Your code must compile using make **without errors**. You will be provided with a Makefile and instructions on using it.
- **Test your program thoroughly with the GCC compiler in a Linux environment.**
- If your code does not compile, **then you will score zero**. Therefore, ensure you have removed all syntax errors from your code.
- **Gradescope** platform would be used to upload the assignment file(s) for grading. The link to the [Gradescope assignment](#) is available on Myls course page. For submission, Drag and drop your code file(s) into Gradescope. **Make sure your file name is as suggested in the assignment; using a different name may score Zero.**
- Please note that the submitted code will be checked for plagiarism. By submitting the code file(s), you would confirm that you have not received unauthorized assistance in preparing the assignment. You also confirm that you are aware of course policies for submitted work.
- Marks will be deducted for any questions where these requirements are not met.
- Multiple attempts will be allowed, but only your last submission before the deadline will be graded. We reserve the right to take off points for not following directions.

## Question 1

Write a C program to validate the solution to the *Sudoku* puzzle. A *Sudoku* puzzle uses a  $9 \times 9$  grid in which each column and row, as well as each of the nine  $3 \times 3$  sub-grids, must contain all the digits  $1 \cdots 9$ . Figure below presents an example of a valid  $9 \times 9$  Sudoku puzzle solution. This program implements a multithreaded application that reads a Sudoku puzzle solution from a *file* ("*sample\_in\_sudoku.txt*") and validates it.

2	7	6	3	1	4	9	5	8
8	5	4	9	6	2	7	1	3
9	1	3	8	7	5	2	6	4
4	6	8	1	2	7	3	9	5
5	9	7	4	3	8	6	2	1
1	3	2	5	9	6	4	8	7
3	2	5	7	8	9	1	4	6
6	4	1	2	5	3	8	7	9
7	8	9	6	4	1	5	3	2

There are several different ways of multi-threading this application (students are encouraged to explore). One suggested strategy is to create threads that check the following criteria:

- A thread to check that each column contains the digits 1 through 9
- A thread to check that each row contains the digits 1 through 9
- Nine threads to check that each of the  $3 \times 3$  sub-grids contain the digits 1 through 9

This would result in a total of eleven separate threads for validating a Sudoku puzzle. However, you are welcome to create even more threads for this problem. For example, rather than creating one thread that checks all nine columns, you could create nine separate threads and have each of them check one column.

1. **Passing Parameters to Each Thread:** The parent thread will create the worker threads, passing each worker the location that it must check in the Sudoku grid. This step will require passing several parameters to each thread. The easiest approach is to create a data structure using a struct. For example, a structure to pass the row and column where a thread must begin validating would appear as follows:

```
/* structure for passing data to threads */
typedef struct
{
    int row;
    int column;
} parameters;
```

Pthreads will create worker threads using a strategy similar to that shown below:

```
parameters *data = (parameters *) malloc(sizeof(parameters));
data->row = 1;
data->column = 1;
/* Now create the thread passing it data as a parameter */
```

The data pointer will be passed to the pthread's `create()` (Pthreads) function, which in turn will pass it as a parameter to the function that is to run as a separate thread.

2. **Returning Results to the Parent Thread:** Each worker thread is assigned the task of determining the validity of a particular region of the Sudoku puzzle. Once a worker has performed this check, it must pass its results back to the parent. One good way to handle this is to create an array of integer values that is visible to each thread. The  $i^{th}$  index in this array corresponds to the  $i^{th}$  worker thread. If a worker sets its corresponding value to 1, it is indicating that its region of the Sudoku puzzle is valid. A value of 0 indicates otherwise. When all worker threads have completed, the parent thread checks each entry in the result array to determine if the Sudoku puzzle is valid.

The expected output for this question is:

```
Sudoku Puzzle Solution is:
2 7 6 3 1 4 9 5 8
8 5 4 9 6 2 7 1 3
9 1 3 8 7 5 2 6 4
4 6 8 1 2 7 3 9 5
5 9 7 4 3 8 6 2 1
1 3 2 5 9 6 4 8 7
3 2 5 7 8 9 1 4 6
6 4 1 2 5 3 8 7 9
7 8 9 6 4 1 5 3 2
Sudoku puzzle is valid
```

**Note:** When submitting source code file for this question, name it like:

- `sudoku.c`

## Question 2

Write a C program to find the average waiting time and turn-around time for the pre-defined set of tasks. The program will read a predefined set of tasks from file “sample\_in\_schedule.txt” and will schedule the tasks based on the First Come First Served scheduling algorithm (non preemptive). The columns in “sample\_in\_schedule.txt” present the thread id, its arrival time, and its burst time, respectively, with the following example format:

1,	0,	3
2,	1,	2
3,	2,	1
4,	3,	4

Thus, thread T2 has arrived at 0 milliseconds and a CPU burst of 3 milliseconds, and so forth. There are several different ways of developing this program (students are encouraged to explore). You can use array's to hold the information and compute the results. But the easiest approach may be to create a data structure using a struct that may appear as follows:

```
struct threadInfo{
    int p_id;
    int arr_time;
    int burst_time;
    int waiting_time;
    int turn_around_time;};
```

Once the arrival time and burst times of threads have been read from the file, the turn-around time and the threads' waiting time are calculated using the following formula.

- $\text{Turn-around\_time} = \text{thread\_completion\_time} - \text{thread\_arrival\_time}$
- $\text{Waiting\_time} = \text{Turn-around\_time} - \text{Thread\_burst\_time}$

The average waiting time for all threads is determined by summing all the threads' respective waiting times and dividing the sum by the total number of threads. Similarly, the average turn-around time is calculated summing the respective turn-around time of all the threads and divided the sum by the total number of threads.

The expected output for this question is:

Thread ID	Arrival Time	Burst Time	Completion Time	Turn-Around Time	Waiting Time
1	0	2	2	2	0
2	1	6	8	7	1
3	2	4	12	10	6
4	3	9	21	18	9
5	6	12	33	27	15
The average waiting time: 6.20					
The average turn-around time: 12.80					

**Note:** When submitting source code file for this question, name it like:

- `fcfs.c`